# IMGPROC. IMAGE PROCESSING

## 3.1 Image Filtering

Functions and classes described in this section are used to perform various linear or non-linear filtering operations on 2D images (represented as Mat()'s). It means that for each pixel location $(x, y)$ in the source image (normally, rectangular), its neighborhood is considered and used to compute the response. In case of a linear filter, it is a weighted sum of pixel values. In case of morphological operations, it is the minimum or maximum values, and so on. The computed response is stored in the destination image at the same location $(x, y)$ . It means that the output image will be of the same size as the input image. Normally, the functions support multi-channel arrays, in which case every channel is processed independently. Therefore, the output image will also have the same number of channels as the input one.

Another common feature of the functions and classes described in this section is that, unlike simple arithmetic functions, they need to extrapolate values of some non-existing pixels. For example, if you want to smooth an image using a Gaussian $3 \times 3$ filter, then, when processing the left-most pixels in each row, you need pixels to the left of them, that is, outside of the image. You can let these pixels be the same as the left-most image pixels ("replicated border" extrapolation method), or assume that all the non-existing pixels are zeros ("constant border" extrapolation method), and so on. OpenCV enables you to specify the extrapolation method. For details, see the function borderInterpolate() and discussion of the borderType parameter in the section and various functions below.

```
/*
 Various border types, image boundaries are denoted with '|'

 * BORDER_REPLICATE:     aaaaaa|abcdefgh|hhhhhhh
 * BORDER_REFLECT:       fedcba|abcdefgh|hgfedcb
 * BORDER_REFLECT_101:   gfedcb|abcdefgh|gfedcba
 * BORDER_WRAP:          cdefgh|abcdefgh|abcdefg
 * BORDER_CONSTANT:      iiiiii|abcdefgh|iiiiiii  with some specified 'i'
 */
```

### BaseColumnFilter

**class BaseColumnFilter**

Base class for filters with single-column kernels.

```
class BaseColumnFilter
{
public:
    virtual ~BaseColumnFilter();

    // To be overriden by the user.
```

```
    //
    // runs a filtering operation on the set of rows,
    // "dstcount + ksize - 1" rows on input,
    // "dstcount" rows on output,
    // each input and output row has "width" elements
    // the filtered rows are written into "dst" buffer.
    virtual void operator()(const uchar** src, uchar* dst, int dststep,
                            int dstcount, int width) = 0;
    // resets the filter state (may be needed for IIR filters)
    virtual void reset();

    int ksize; // the aperture size
    int anchor; // position of the anchor point,
                // normally not used during the processing
};
```

The class `BaseColumnFilter` is a base class for filtering data using single-column kernels. Filtering does not have to be a linear operation. In general, it could be written as follows:

$$dst(x, y) = F(src[y](x), \ src[y+1](x), \ ..., \ src[y+ksize-1](x))$$

where $F$ is a filtering function but, as it is represented as a class, it can produce any side effects, memorize previously processed data, and so on. The class only defines an interface and is not used directly. Instead, there are several functions in OpenCV (and you can add more) that return pointers to the derived classes that implement specific filtering operations. Those pointers are then passed to the `FilterEngine` constructor. While the filtering operation interface uses the `uchar` type, a particular implementation is not limited to 8-bit data.

**See Also:**

BaseRowFilter, BaseFilter, FilterEngine, getColumnSumFilter(), getLinearColumnFilter(), getMorphologyColumnFilter()

## BaseFilter

**class BaseFilter**

Base class for 2D image filters.

```
class BaseFilter
{
public:
    virtual ~BaseFilter();

    // To be overriden by the user.
    //
    // runs a filtering operation on the set of rows,
    // "dstcount + ksize.height - 1" rows on input,
    // "dstcount" rows on output,
    // each input row has "(width + ksize.width-1)*cn" elements
    // each output row has "width*cn" elements.
    // the filtered rows are written into "dst" buffer.
    virtual void operator()(const uchar** src, uchar* dst, int dststep,
                            int dstcount, int width, int cn) = 0;
    // resets the filter state (may be needed for IIR filters)
    virtual void reset();
    Size ksize;
    Point anchor;
};
```

The class `BaseFilter` is a base class for filtering data using 2D kernels. Filtering does not have to be a linear operation. In general, it could be written as follows:

$$\mathtt{dst}(x,y) = F(\mathtt{src}[y](x),\ \mathtt{src}[y](x+1),\ ...,\ \mathtt{src}[y](x+\mathtt{ksize.width}-1),$$
$$\mathtt{src}[y+1](x),\ \mathtt{src}[y+1](x+1),\ ...,\ \mathtt{src}[y+1](x+\mathtt{ksize.width}-1),$$
$$\text{.......................................................................................}$$
$$\mathtt{src}[y+\mathtt{ksize.height\text{-}1}](x),$$
$$\mathtt{src}[y+\mathtt{ksize.height\text{-}1}](x+1),$$
$$...\mathtt{src}[y+\mathtt{ksize.height\text{-}1}](x+\mathtt{ksize.width}-1))$$

where $F$ is a filtering function. The class only defines an interface and is not used directly. Instead, there are several functions in OpenCV (and you can add more) that return pointers to the derived classes that implement specific filtering operations. Those pointers are then passed to the `FilterEngine` constructor. While the filtering operation interface uses the `uchar` type, a particular implementation is not limited to 8-bit data.

**See Also:**

BaseColumnFilter, BaseRowFilter, FilterEngine, getLinearFilter(), getMorphologyFilter()

## BaseRowFilter

**class BaseRowFilter**

Base class for filters with single-row kernels.

```cpp
class BaseRowFilter
{
public:
    virtual ~BaseRowFilter();

    // To be overriden by the user.
    //
    // runs filtering operation on the single input row
    // of "width" element, each element is has "cn" channels.
    // the filtered row is written into "dst" buffer.
    virtual void operator()(const uchar* src, uchar* dst,
                            int width, int cn) = 0;
    int ksize, anchor;
};
```

The class `BaseRowFilter` is a base class for filtering data using single-row kernels. Filtering does not have to be a linear operation. In general, it could be written as follows:

$$\mathtt{dst}(x,y) = F(\mathtt{src}[y](x),\ \mathtt{src}[y](x+1),\ ...,\ \mathtt{src}[y](x+\mathtt{ksize.width}-1))$$

where $F$ is a filtering function. The class only defines an interface and is not used directly. Instead, there are several functions in OpenCV (and you can add more) that return pointers to the derived classes that implement specific filtering operations. Those pointers are then passed to the `FilterEngine` constructor. While the filtering operation interface uses the `uchar` type, a particular implementation is not limited to 8-bit data.

**See Also:**

BaseColumnFilter, BaseFilter, FilterEngine, getLinearRowFilter(), getMorphologyRowFilter(), getRowSumFilter()

## FilterEngine

**class FilterEngine**

Generic image filtering class.

```
class FilterEngine
{
public:
    // empty constructor
    FilterEngine();
    // builds a 2D non-separable filter (!_filter2D.empty()) or
    // a separable filter (!_rowFilter.empty() && !_columnFilter.empty())
    // the input data type will be "srcType", the output data type will be "dstType",
    // the intermediate data type is "bufType".
    // _rowBorderType and _columnBorderType determine how the image
    // will be extrapolated beyond the image boundaries.
    // _borderValue is only used when _rowBorderType and/or _columnBorderType
    // == BORDER_CONSTANT
    FilterEngine(const Ptr<BaseFilter>& _filter2D,
                 const Ptr<BaseRowFilter>& _rowFilter,
                 const Ptr<BaseColumnFilter>& _columnFilter,
                 int srcType, int dstType, int bufType,
                 int _rowBorderType=BORDER_REPLICATE,
                 int _columnBorderType=-1, // use _rowBorderType by default
                 const Scalar& _borderValue=Scalar());
    virtual ~FilterEngine();
    // separate function for the engine initialization
    void init(const Ptr<BaseFilter>& _filter2D,
            const Ptr<BaseRowFilter>& _rowFilter,
            const Ptr<BaseColumnFilter>& _columnFilter,
            int srcType, int dstType, int bufType,
            int _rowBorderType=BORDER_REPLICATE, int _columnBorderType=-1,
            const Scalar& _borderValue=Scalar());
    // starts filtering of the ROI in an image of size "wholeSize".
    // returns the starting y-position in the source image.
    virtual int start(Size wholeSize, Rect roi, int maxBufRows=-1);
    // alternative form of start that takes the image
    // itself instead of "wholeSize". Set isolated to true to pretend that
    // there are no real pixels outside of the ROI
    // (so that the pixels are extrapolated using the specified border modes)
    virtual int start(const Mat& src, const Rect& srcRoi=Rect(0,0,-1,-1),
                      bool isolated=false, int maxBufRows=-1);
    // processes the next portion of the source image,
    // "srcCount" rows starting from "src" and
    // stores the results in "dst".
    // returns the number of produced rows
    virtual int proceed(const uchar* src, int srcStep, int srcCount,
                        uchar* dst, int dstStep);
    // higher-level function that processes the whole
    // ROI or the whole image with a single call
    virtual void apply( const Mat& src, Mat& dst,
                        const Rect& srcRoi=Rect(0,0,-1,-1),
                        Point dstOfs=Point(0,0),
                        bool isolated=false);
    bool isSeparable() const { return filter2D.empty(); }
    // how many rows from the input image are not yet processed
    int remainingInputRows() const;
    // how many output rows are not yet produced
    int remainingOutputRows() const;
    ...
    // the starting and the ending rows in the source image
    int startY, endY;
```

```
    // pointers to the filters
    Ptr<BaseFilter> filter2D;
    Ptr<BaseRowFilter> rowFilter;
    Ptr<BaseColumnFilter> columnFilter;
};
```

The class `FilterEngine` can be used to apply an arbitrary filtering operation to an image. It contains all the necessary intermediate buffers, computes extrapolated values of the "virtual" pixels outside of the image, and so on. Pointers to the initialized `FilterEngine` instances are returned by various `create*Filter` functions (see below) and they are used inside high-level functions such as `filter2D()`, `erode()`, `dilate()`, and others. Thus, the class plays a key role in many of OpenCV filtering functions.

This class makes it easier to combine filtering operations with other operations, such as color space conversions, thresholding, arithmetic operations, and others. By combining several operations together you can get much better performance because your data will stay in cache. For example, see below the implementation of the Laplace operator for floating-point images, which is a simplified implementation of `Laplacian()` :

```
void laplace_f(const Mat& src, Mat& dst)
{
    CV_Assert( src.type() == CV_32F );
    dst.create(src.size(), src.type());

    // get the derivative and smooth kernels for d2I/dx2.
    // for d2I/dy2 consider using the same kernels, just swapped
    Mat kd, ks;
    getSobelKernels( kd, ks, 2, 0, ksize, false, ktype );

    // process 10 source rows at once
    int DELTA = std::min(10, src.rows);
    Ptr<FilterEngine> Fxx = createSeparableLinearFilter(src.type(),
        dst.type(), kd, ks, Point(-1,-1), 0, borderType, borderType, Scalar() );
    Ptr<FilterEngine> Fyy = createSeparableLinearFilter(src.type(),
        dst.type(), ks, kd, Point(-1,-1), 0, borderType, borderType, Scalar() );

    int y = Fxx->start(src), dsty = 0, dy = 0;
    Fyy->start(src);
    const uchar* sptr = src.data + y*src.step;

    // allocate the buffers for the spatial image derivatives;
    // the buffers need to have more than DELTA rows, because at the
    // last iteration the output may take max(kd.rows-1,ks.rows-1)
    // rows more than the input.
    Mat Ixx( DELTA + kd.rows - 1, src.cols, dst.type() );
    Mat Iyy( DELTA + kd.rows - 1, src.cols, dst.type() );

    // inside the loop always pass DELTA rows to the filter
    // (note that the "proceed" method takes care of possibe overflow, since
    // it was given the actual image height in the "start" method)
    // on output you can get:
    //  * < DELTA rows (initial buffer accumulation stage)
    //  * = DELTA rows (settled state in the middle)
    //  * > DELTA rows (when the input image is over, generate
    //                  "virtual" rows using the border mode and filter them)
    // this variable number of output rows is dy.
    // dsty is the current output row.
    // sptr is the pointer to the first input row in the portion to process
    for( ; dsty < dst.rows; sptr += DELTA*src.step, dsty += dy )
    {
```

```
        Fxx->proceed( sptr, (int)src.step, DELTA, Ixx.data, (int)Ixx.step );
        dy = Fyy->proceed( sptr, (int)src.step, DELTA, d2y.data, (int)Iyy.step );
        if( dy > 0 )
        {
            Mat dstripe = dst.rowRange(dsty, dsty + dy);
            add(Ixx.rowRange(0, dy), Iyy.rowRange(0, dy), dstripe);
        }
    }
}
```

If you do not need that much control of the filtering process, you can simply use the `FilterEngine::apply` method. The method is implemented as follows:

```
void FilterEngine::apply(const Mat& src, Mat& dst,
    const Rect& srcRoi, Point dstOfs, bool isolated)
{
    // check matrix types
    CV_Assert( src.type() == srcType && dst.type() == dstType );

    // handle the "whole image" case
    Rect _srcRoi = srcRoi;
    if( _srcRoi == Rect(0,0,-1,-1) )
        _srcRoi = Rect(0,0,src.cols,src.rows);

    // check if the destination ROI is inside dst.
    // and FilterEngine::start will check if the source ROI is inside src.
    CV_Assert( dstOfs.x >= 0 && dstOfs.y >= 0 &&
        dstOfs.x + _srcRoi.width <= dst.cols &&
        dstOfs.y + _srcRoi.height <= dst.rows );

    // start filtering
    int y = start(src, _srcRoi, isolated);

    // process the whole ROI. Note that "endY - startY" is the total number
    // of the source rows to process
    // (including the possible rows outside of srcRoi but inside the source image)
    proceed( src.data + y*src.step,
            (int)src.step, endY - startY,
            dst.data + dstOfs.y*dst.step +
            dstOfs.x*dst.elemSize(), (int)dst.step );
}
```

Unlike the earlier versions of OpenCV, now the filtering operations fully support the notion of image ROI, that is, pixels outside of the ROI but inside the image can be used in the filtering operations. For example, you can take a ROI of a single pixel and filter it. This will be a filter response at that particular pixel. However, it is possible to emulate the old behavior by passing `isolated=false` to `FilterEngine::start` or `FilterEngine::apply`. You can pass the ROI explicitly to `FilterEngine::apply` or construct new matrix headers:

```
// compute dI/dx derivative at src(x,y)

// method 1:
// form a matrix header for a single value
float val1 = 0;
Mat dst1(1,1,CV_32F,&val1);

Ptr<FilterEngine> Fx = createDerivFilter(CV_32F, CV_32F,
                    1, 0, 3, BORDER_REFLECT_101);
Fx->apply(src, Rect(x,y,1,1), Point(), dst1);
```

```
// method 2:
// form a matrix header for a single value
float val2 = 0;
Mat dst2(1,1,CV_32F,&val2);

Mat pix_roi(src, Rect(x,y,1,1));
Sobel(pix_roi, dst2, dst2.type(), 1, 0, 3, 1, 0, BORDER_REFLECT_101);

printf("method1 =
```

Explore the data types. As it was mentioned in the `BaseFilter` description, the specific filters can process data of any type, despite that `Base*Filter::operator()` only takes `uchar` pointers and no information about the actual types. To make it all work, the following rules are used:

- In case of separable filtering, `FilterEngine::rowFilter` is applied first. It transforms the input image data (of type `srcType` ) to the intermediate results stored in the internal buffers (of type `bufType` ). Then, these intermediate results are processed as *single-channel data* with `FilterEngine::columnFilter` and stored in the output image (of type `dstType` ). Thus, the input type for `rowFilter` is `srcType` and the output type is `bufType` . The input type for `columnFilter` is `CV_MAT_DEPTH(bufType)` and the output type is `CV_MAT_DEPTH(dstType)` .

- In case of non-separable filtering, `bufType` must be the same as `srcType` . The source data is copied to the temporary buffer, if needed, and then just passed to `FilterEngine::filter2D` . That is, the input type for `filter2D` is `srcType` (= `bufType` ) and the output type is `dstType` .

**See Also:**

BaseColumnFilter, BaseFilter, BaseRowFilter, createBoxFilter(), createDerivFilter(), createGaussianFilter(), createLinearFilter(), createMorphologyFilter(), createSeparableLinearFilter()

## bilateralFilter

Applies the bilateral filter to an image.

**C++:** void **bilateralFilter**(InputArray **src**, OutputArray **dst**, int **d**, double **sigmaColor**, double **sigmaSpace**, int **borderType**=BORDER_DEFAULT )

**Python:** cv2.**bilateralFilter**(src, d, sigmaColor, sigmaSpace[, dst[, borderType ]]) → dst

> **Parameters**

> > **src** – Source 8-bit or floating-point, 1-channel or 3-channel image.

> > **dst** – Destination image of the same size and type as `src` .

> > **d** – Diameter of each pixel neighborhood that is used during filtering. If it is non-positive, it is computed from `sigmaSpace` .

> > **sigmaColor** – Filter sigma in the color space. A larger value of the parameter means that farther colors within the pixel neighborhood (see `sigmaSpace` ) will be mixed together, resulting in larger areas of semi-equal color.

> > **sigmaSpace** – Filter sigma in the coordinate space. A larger value of the parameter means that farther pixels will influence each other as long as their colors are close enough (see `sigmaColor` ). When `d>0` , it specifies the neighborhood size regardless of `sigmaSpace` . Otherwise, `d` is proportional to `sigmaSpace` .

The function applies bilateral filtering to the input image, as described in http://www.dai.ed.ac.uk/CVonline/LOCAL_COPIES/MANDUCHI1/Bilateral_Filtering.html `bilateralFilter`

can reduce unwanted noise very well while keeping edges fairly sharp. However, it is very slow compared to most filters.

*Sigma values*: For simplicity, you can set the 2 sigma values to be the same. If they are small (< 10), the filter will not have much effect, whereas if they are large (> 150), they will have a very strong effect, making the image look "cartoonish".

*Filter size*: Large filters (d > 5) are very slow, so it is recommended to use d=5 for real-time applications, and perhaps d=9 for offline applications that need heavy noise filtering.

This filter does not work inplace.

## blur

Smoothes an image using the normalized box filter.

**C++:** void **blur**(InputArray **src**, OutputArray **dst**, Size **ksize**, Point **anchor**=Point(-1,-1), int **border-Type**=BORDER_DEFAULT )

**Python:** cv2.**blur**(src, ksize[, dst[, anchor[, borderType ]]]) → dst

>   **Parameters**
>
>>   **src** – Source image. The image can have any number of channels, which are processed independently. The depth should be `CV_8U`, `CV_16U`, `CV_16S`, `CV_32F` or `CV_64F`.
>>
>>   **dst** – Destination image of the same size and type as `src` .
>>
>>   **ksize** – Smoothing kernel size.
>>
>>   **anchor** – Anchor point. The default value `Point(-1,-1)` means that the anchor is at the kernel center.
>>
>>   **borderType** – Border mode used to extrapolate pixels outside of the image.

The function smoothes an image using the kernel:

$$K = \frac{1}{\texttt{ksize.width} * \texttt{ksize.height}} \begin{bmatrix} 1 & 1 & 1 & \cdots & 1 & 1 \\ 1 & 1 & 1 & \cdots & 1 & 1 \\ & & \cdots\cdots\cdots & & \\ 1 & 1 & 1 & \cdots & 1 & 1 \end{bmatrix}$$

The call `blur(src, dst, ksize, anchor, borderType)` is equivalent to `boxFilter(src, dst, src.type(), anchor, true, borderType)` .

**See Also:**

`boxFilter()`, `bilateralFilter()`, `GaussianBlur()`, `medianBlur()`

## borderInterpolate

Computes the source location of an extrapolated pixel.

**C++:** int **borderInterpolate**(int **p**, int **len**, int **borderType**)

**Python:** cv2.**borderInterpolate**(p, len, borderType) → retval

>   **Parameters**
>
>>   **p** – 0-based coordinate of the extrapolated pixel along one of the axes, likely <0 or >= `len` .
>>
>>   **len** – Length of the array along the corresponding axis.

> **borderType** – Border type, one of the BORDER_* , except for BORDER_TRANSPARENT and BORDER_ISOLATED . When borderType==BORDER_CONSTANT , the function always returns -1, regardless of p and len .

The function computes and returns the coordinate of a donor pixel corresponding to the specified extrapolated pixel when using the specified extrapolation border mode. For example, if you use BORDER_WRAP mode in the horizontal direction, BORDER_REFLECT_101 in the vertical direction and want to compute value of the "virtual" pixel Point(-5, 100) in a floating-point image img , it looks like:

```
float val = img.at<float>(borderInterpolate(100, img.rows, BORDER_REFLECT_101),
                          borderInterpolate(-5, img.cols, BORDER_WRAP));
```

Normally, the function is not called directly. It is used inside FilterEngine and copyMakeBorder() to compute tables for quick extrapolation.

**See Also:**

FilterEngine, copyMakeBorder()

## boxFilter

Smoothes an image using the box filter.

**C++:** void **boxFilter**(InputArray **src**, OutputArray **dst**, int **ddepth**, Size **ksize**, Point **anchor**=Point(-1,-1), bool **normalize**=true, int **borderType**=BORDER_DEFAULT )

**Python:** cv2.**boxFilter**(src, ddepth, ksize[, dst[, anchor[, normalize[, borderType ]]]]) → dst

> **Parameters**
>
> > **src** – Source image.
> >
> > **dst** – Destination image of the same size and type as src .
> >
> > **ksize** – Smoothing kernel size.
> >
> > **anchor** – Anchor point. The default value Point(-1,-1) means that the anchor is at the kernel center.
> >
> > **normalize** – Flag specifying whether the kernel is normalized by its area or not.
> >
> > **borderType** – Border mode used to extrapolate pixels outside of the image.

The function smoothes an image using the kernel:

$$
K = \alpha \begin{bmatrix} 1 & 1 & 1 & \cdots & 1 & 1 \\ 1 & 1 & 1 & \cdots & 1 & 1 \\ & & \cdots\cdots\cdots & & \\ 1 & 1 & 1 & \cdots & 1 & 1 \end{bmatrix}
$$

where

$$
\alpha = \begin{cases} \frac{1}{\texttt{ksize.width}*\texttt{ksize.height}} & \text{when normalize=true} \\ 1 & \text{otherwise} \end{cases}
$$

Unnormalized box filter is useful for computing various integral characteristics over each pixel neighborhood, such as covariance matrices of image derivatives (used in dense optical flow algorithms, and so on). If you need to compute pixel sums over variable-size windows, use integral() .

**See Also:**

boxFilter(), bilateralFilter(), GaussianBlur(), medianBlur(), integral()

---

## buildPyramid

Constructs the Gaussian pyramid for an image.

**C++:** void **buildPyramid**(InputArray **src**, OutputArrayOfArrays **dst**, int **maxlevel**)

>> **Parameters**

>>> **src** – Source image. Check `pyrDown()` for the list of supported types.

>>> **dst** – Destination vector of `maxlevel+1` images of the same type as `src` . `dst[0]` will be the same as `src` . `dst[1]` is the next pyramid layer, a smoothed and down-sized `src` , and so on.

>>> **maxlevel** – 0-based index of the last (the smallest) pyramid layer. It must be non-negative.

The function constructs a vector of images and builds the Gaussian pyramid by recursively applying `pyrDown()` to the previously built pyramid layers, starting from `dst[0]==src` .

## copyMakeBorder

Forms a border around an image.

**C++:** void **copyMakeBorder**(InputArray **src**, OutputArray **dst**, int **top**, int **bottom**, int **left**, int **right**, int **borderType**, const Scalar& **value**=Scalar() )

**Python:** cv2.**copyMakeBorder**(src, top, bottom, left, right, borderType[, dst[, value ] ]) → dst

**C:** void **cvCopyMakeBorder**(const CvArr* **src**, CvArr* **dst**, CvPoint **offset**, int **bordertype**, CvScalar **value**=cvScalarAll(0) )

**Python:** cv.**CopyMakeBorder**(src, dst, offset, bordertype, value=(0, 0, 0, 0)) → None

>> **Parameters**

>>> **src** – Source image.

>>> **dst** – Destination image of the same type as `src` and the size `Size(src.cols+left+right, src.rows+top+bottom)` .

>>> **top** –

>>> **bottom** –

>>> **left** –

>>> **right** – Parameter specifying how many pixels in each direction from the source image rectangle to extrapolate. For example, `top=1, bottom=1, left=1, right=1` mean that 1 pixel-wide border needs to be built.

>>> **borderType** – Border type. See `borderInterpolate()` for details.

>>> **value** – Border value if `borderType==BORDER_CONSTANT` .

The function copies the source image into the middle of the destination image. The areas to the left, to the right, above and below the copied source image will be filled with extrapolated pixels. This is not what `FilterEngine` or filtering functions based on it do (they extrapolate pixels on-fly), but what other more complex functions, including your own, may do to simplify image boundary handling.

The function supports the mode when `src` is already in the middle of `dst` . In this case, the function does not copy `src` itself but simply constructs the border, for example:

```
// let border be the same in all directions
int border=2;
// constructs a larger image to fit both the image and the border
Mat gray_buf(rgb.rows + border*2, rgb.cols + border*2, rgb.depth());
// select the middle part of it w/o copying data
Mat gray(gray_canvas, Rect(border, border, rgb.cols, rgb.rows));
// convert image from RGB to grayscale
cvtColor(rgb, gray, CV_RGB2GRAY);
// form a border in-place
copyMakeBorder(gray, gray_buf, border, border,
               border, border, BORDER_REPLICATE);
// now do some custom filtering ...
...
```

**Note:** When the source image is a part (ROI) of a bigger image, the function will try to use the pixels outside of the ROI to form a border. To disable this feature and always do extrapolation, as if `src` was not a ROI, use `borderType | BORDER_ISOLATED`.

**See Also:**

`borderInterpolate()`

## createBoxFilter

Returns a box filter engine.

**C++:** Ptr<FilterEngine> **createBoxFilter**(int **srcType**, int **dstType**, Size **ksize**, Point **anchor**=Point(-1,-1), bool **normalize**=true, int **borderType**=BORDER_DEFAULT)

**C++:** Ptr<BaseRowFilter> **getRowSumFilter**(int **srcType**, int **sumType**, int **ksize**, int **anchor**=-1)

**C++:** Ptr<BaseColumnFilter> **getColumnSumFilter**(int **sumType**, int **dstType**, int **ksize**, int **anchor**=-1, double **scale**=1)

>    **Parameters**

>>        **srcType** – Source image type.

>>        **sumType** – Intermediate horizontal sum type that must have as many channels as `srcType` .

>>        **dstType** – Destination image type that must have as many channels as `srcType` .

>>        **ksize** – Aperture size.

>>        **anchor** – Anchor position with the kernel. Negative values mean that the anchor is at the kernel center.

>>        **normalize** – Flag specifying whether the sums are normalized or not. See `boxFilter()` for details.

>>        **scale** – Another way to specify normalization in lower-level `getColumnSumFilter` .

>>        **borderType** – Border type to use. See `borderInterpolate()` .

The function is a convenience function that retrieves the horizontal sum primitive filter with `getRowSumFilter()` , vertical sum filter with `getColumnSumFilter()` , constructs new `FilterEngine` , and passes both of the primitive filters there. The constructed filter engine can be used for image filtering with normalized or unnormalized box filter.

The function itself is used by `blur()` and `boxFilter()` .

---

**See Also:**

FilterEngine, blur(), boxFilter()

## createDerivFilter

Returns an engine for computing image derivatives.

**C++:** Ptr<FilterEngine> **createDerivFilter**(int **srcType**, int **dstType**, int **dx**, int **dy**, int **ksize**, int **border-Type**=BORDER_DEFAULT )

> **Parameters**
>
>> **srcType** – Source image type.
>>
>> **dstType** – Destination image type that must have as many channels as srcType .
>>
>> **dx** – Derivative order in respect of x.
>>
>> **dy** – Derivative order in respect of y.
>>
>> **ksize** – Aperture size See getDerivKernels() .
>>
>> **borderType** – Border type to use. See borderInterpolate() .

The function createDerivFilter() is a small convenience function that retrieves linear filter coefficients for computing image derivatives using getDerivKernels() and then creates a separable linear filter with createSeparableLinearFilter() . The function is used by Sobel() and Scharr() .

**See Also:**

createSeparableLinearFilter(), getDerivKernels(), Scharr(), Sobel()

## createGaussianFilter

Returns an engine for smoothing images with the Gaussian filter.

**C++:** Ptr<FilterEngine> **createGaussianFilter**(int **type**, Size **ksize**, double **sigmaX**, double **sigmaY**=0, int **borderType**=BORDER_DEFAULT)

> **Parameters**
>
>> **type** – Source and destination image type.
>>
>> **ksize** – Aperture size. See getGaussianKernel() .
>>
>> **sigmaX** – Gaussian sigma in the horizontal direction. See getGaussianKernel() .
>>
>> **sigmaY** – Gaussian sigma in the vertical direction. If 0, then sigmaY $\leftarrow$ sigmaX .
>>
>> **borderType** – Border type to use. See borderInterpolate() .

The function createGaussianFilter() computes Gaussian kernel coefficients and then returns a separable linear filter for that kernel. The function is used by GaussianBlur() . Note that while the function takes just one data type, both for input and output, you can pass this limitation by calling getGaussianKernel() and then createSeparableLinearFilter() directly.

**See Also:**

createSeparableLinearFilter(), getGaussianKernel(), GaussianBlur()

## createLinearFilter

Creates a non-separable linear filter engine.

**C++:** Ptr<FilterEngine> **createLinearFilter**(int **srcType**, int **dstType**, InputArray **kernel**, Point **_anchor**=Point(-1,-1), double **delta**=0, int **rowBorderType**=BORDER_DEFAULT, int **columnBorderType**=-1, const Scalar& **borderValue**=Scalar())

**C++:** Ptr<BaseFilter> **getLinearFilter**(int **srcType**, int **dstType**, InputArray **kernel**, Point **anchor**=Point(-1,-1), double **delta**=0, int **bits**=0)

> **Parameters**
>
> > **srcType** – Source image type.
> >
> > **dstType** – Destination image type that must have as many channels as `srcType` .
> >
> > **kernel** – 2D array of filter coefficients.
> >
> > **anchor** – Anchor point within the kernel. Special value `Point(-1,-1)` means that the anchor is at the kernel center.
> >
> > **delta** – Value added to the filtered results before storing them.
> >
> > **bits** – Number of the fractional bits. The parameter is used when the kernel is an integer matrix representing fixed-point filter coefficients.
> >
> > **rowBorderType** – Pixel extrapolation method in the vertical direction. For details, see `borderInterpolate()`.
> >
> > **columnBorderType** – Pixel extrapolation method in the horizontal direction.
> >
> > **borderValue** – Border value used in case of a constant border.

The function returns a pointer to a 2D linear filter for the specified kernel, the source array type, and the destination array type. The function is a higher-level function that calls `getLinearFilter` and passes the retrieved 2D filter to the `FilterEngine` constructor.

**See Also:**

`createSeparableLinearFilter()`, `FilterEngine`, `filter2D()`

## createMorphologyFilter

Creates an engine for non-separable morphological operations.

**C++:** Ptr<FilterEngine> **createMorphologyFilter**(int **op**, int **type**, InputArray **element**, Point **anchor**=Point(-1,-1), int **rowBorderType**=BORDER_CONSTANT, int **columnBorderType**=-1, const Scalar& **borderValue**=morphologyDefaultBorderValue())

**C++:** Ptr<BaseFilter> **getMorphologyFilter**(int **op**, int **type**, InputArray **element**, Point **anchor**=Point(-1,-1))

**C++:** Ptr<BaseRowFilter> **getMorphologyRowFilter**(int **op**, int **type**, int **esize**, int **anchor**=-1)

**C++:** Ptr<BaseColumnFilter> **getMorphologyColumnFilter**(int **op**, int **type**, int **esize**, int **anchor**=-1)

**C++:** Scalar **morphologyDefaultBorderValue**()

> **Parameters**
>
> > **op** – Morphology operation ID, `MORPH_ERODE` or `MORPH_DILATE` .

        **type** – Input/output image type. The number of channels can be arbitrary. The depth should be one of `CV_8U`, `CV_16U`, `CV_16S`, `CV_32F'` or `''CV_64F`.

        **element** – 2D 8-bit structuring element for a morphological operation. Non-zero elements indicate the pixels that belong to the element.

        **esize** – Horizontal or vertical structuring element size for separable morphological operations.

        **anchor** – Anchor position within the structuring element. Negative values mean that the anchor is at the kernel center.

        **rowBorderType** – Pixel extrapolation method in the vertical direction. For details, see `borderInterpolate()`.

        **columnBorderType** – Pixel extrapolation method in the horizontal direction.

        **borderValue** – Border value in case of a constant border. The default value, `morphologyDefaultBorderValue`, has a special meaning. It is transformed $+\inf$ for the erosion and to $-\inf$ for the dilation, which means that the minimum (maximum) is effectively computed only over the pixels that are inside the image.

The functions construct primitive morphological filtering operations or a filter engine based on them. Normally it is enough to use `createMorphologyFilter()` or even higher-level `erode()`, `dilate()`, or `morphologyEx()`. Note that `createMorphologyFilter()` analyzes the structuring element shape and builds a separable morphological filter engine when the structuring element is square.

**See Also:**

`erode()`, `dilate()`, `morphologyEx()`, `FilterEngine`

## createSeparableLinearFilter

Creates an engine for a separable linear filter.

**C++:** Ptr<FilterEngine> **createSeparableLinearFilter**(int **srcType**, int **dstType**, InputArray **rowKernel**, InputArray **columnKernel**, Point **anchor**=Point(-1,-1), double **delta**=0, int **rowBorderType**=BORDER_DEFAULT, int **columnBorderType**=-1, const Scalar& **borderValue**=Scalar())

**C++:** Ptr<BaseColumnFilter> **getLinearColumnFilter**(int **bufType**, int **dstType**, InputArray **columnKernel**, int **anchor**, int **symmetryType**, double **delta**=0, int **bits**=0)

**C++:** Ptr<BaseRowFilter> **getLinearRowFilter**(int **srcType**, int **bufType**, InputArray **rowKernel**, int **anchor**, int **symmetryType**)

    **Parameters**

        **srcType** – Source array type.

        **dstType** – Destination image type that must have as many channels as `srcType`.

        **bufType** – Intermediate buffer type that must have as many channels as `srcType`.

        **rowKernel** – Coefficients for filtering each row.

        **columnKernel** – Coefficients for filtering each column.

        **anchor** – Anchor position within the kernel. Negative values mean that anchor is positioned at the aperture center.

**delta** – Value added to the filtered results before storing them.

**bits** – Number of the fractional bits. The parameter is used when the kernel is an integer matrix representing fixed-point filter coefficients.

**rowBorderType** – Pixel extrapolation method in the vertical direction. For details, see `borderInterpolate()`.

**columnBorderType** – Pixel extrapolation method in the horizontal direction.

**borderValue** – Border value used in case of a constant border.

**symmetryType** – Type of each row and column kernel. See `getKernelType()` .

The functions construct primitive separable linear filtering operations or a filter engine based on them. Normally it is enough to use `createSeparableLinearFilter()` or even higher-level `sepFilter2D()` . The function `createMorphologyFilter()` is smart enough to figure out the `symmetryType` for each of the two kernels, the intermediate `bufType` and, if filtering can be done in integer arithmetics, the number of `bits` to encode the filter coefficients. If it does not work for you, it is possible to call `getLinearColumnFilter`,''getLinearRowFilter'' directly and then pass them to the `FilterEngine` constructor.

**See Also:**

`sepFilter2D()`, `createLinearFilter()`, `FilterEngine`, `getKernelType()`

## dilate

Dilates an image by using a specific structuring element.

**C++:** void **dilate**(InputArray **src**, OutputArray **dst**, InputArray **element**, Point **anchor**=Point(-1,-1), int **iterations**=1, int **borderType**=BORDER_CONSTANT, const Scalar& **border-Value**=morphologyDefaultBorderValue() )

**Python:** cv2.**dilate**(src, kernel[, dst[, anchor[, iterations[, borderType[, borderValue ]]]]]) → dst

**C:** void **cvDilate**(const CvArr* **src**, CvArr* **dst**, IplConvKernel* **element**=NULL, int **iterations**=1 )

**Python:** cv.**Dilate**(src, dst, element=None, iterations=1) → None

> **Parameters**

> > **src** – Source image. The number of channels can be arbitrary. The depth should be one of `CV_8U, CV_16U, CV_16S, CV_32F` or ``CV_64F`.

> > **dst** – Destination image of the same size and type as `src` .

> > **element** – Structuring element used for dilation. If `element=Mat()` , a `3 x 3` rectangular structuring element is used.

> > **anchor** – Position of the anchor within the element. The default value (`-1, -1`) means that the anchor is at the element center.

> > **iterations** – Number of times dilation is applied.

> > **borderType** – Pixel extrapolation method. See `borderInterpolate()` for details.

> > **borderValue** – Border value in case of a constant border. The default value has a special meaning. See `createMorphologyFilter()` for details.

The function dilates the source image using the specified structuring element that determines the shape of a pixel neighborhood over which the maximum is taken:

$$\mathtt{dst}(x, y) = \max_{(x', y'): \, \mathtt{element}(x', y') \neq 0} \mathtt{src}(x + x', y + y')$$

The function supports the in-place mode. Dilation can be applied several ( `iterations` ) times. In case of multi-channel images, each channel is processed independently.

**See Also:**

erode(), morphologyEx(), createMorphologyFilter()

## erode

Erodes an image by using a specific structuring element.

**C++:** void **erode**(InputArray **src**, OutputArray **dst**, InputArray **element**, Point **anchor**=Point(-1,-1), int **iterations**=1, int **borderType**=BORDER_CONSTANT, const Scalar& **borderValue**=morphologyDefaultBorderValue() )

**Python:** cv2.**erode**(src, kernel[, dst[, anchor[, iterations[, borderType[, borderValue ]]]]]) → dst

**C:** void **cvErode**(const CvArr* **src**, CvArr* **dst**, IplConvKernel* **element**=NULL, int **iterations**=1)

**Python:** cv.**Erode**(src, dst, element=None, iterations=1) → None

> **Parameters**
>
> > **src** – Source image. The number of channels can be arbitrary. The depth should be one of `CV_8U, CV_16U, CV_16S, CV_32F‘ or ‘‘CV_64F`.
> >
> > **dst** – Destination image of the same size and type as `src`.
> >
> > **element** – Structuring element used for erosion. If `element=Mat()` , a 3 x 3 rectangular structuring element is used.
> >
> > **anchor** – Position of the anchor within the element. The default value (`-1, -1`) means that the anchor is at the element center.
> >
> > **iterations** – Number of times erosion is applied.
> >
> > **borderType** – Pixel extrapolation method. See borderInterpolate() for details.
> >
> > **borderValue** – Border value in case of a constant border. The default value has a special meaning. See createMorphologyFilter() for details.

The function erodes the source image using the specified structuring element that determines the shape of a pixel neighborhood over which the minimum is taken:

$$\mathtt{dst}(x,y) = \min_{(x',y'):\, \mathtt{element}(x',y') \neq 0} \mathtt{src}(x + x', y + y')$$

The function supports the in-place mode. Erosion can be applied several ( `iterations` ) times. In case of multi-channel images, each channel is processed independently.

**See Also:**

dilate(), morphologyEx(), createMorphologyFilter()

## filter2D

Convolves an image with the kernel.

**C++:** void **filter2D**(InputArray **src**, OutputArray **dst**, int **ddepth**, InputArray **kernel**, Point **anchor**=Point(-1,-1), double **delta**=0, int **borderType**=BORDER_DEFAULT )

**Python:** cv2.**filter2D**(src, ddepth, kernel[, dst[, anchor[, delta[, borderType ]]]]) → dst

**C:** void **cvFilter2D**(const CvArr* **src**, CvArr* **dst**, const CvMat* **kernel**, CvPoint **anchor**=cvPoint(-1, -1))

**Python:** cv.**Filter2D**(src, dst, kernel, anchor=(-1, -1)) → None

>> **Parameters**

>>> **src** – Source image.

>>> **dst** – Destination image of the same size and the same number of channels as `src` .

>>> **ddepth** –

>>> **Desired depth of the destination image. If it is negative, it will be the same as `src.depth()`. The following c**

>>>> – `src.depth()` = CV_8U, ddepth = -1/CV_16S/CV_32F/CV_64F

>>>> – `src.depth()` = CV_16U/CV_16S, ddepth = -1/CV_32F/CV_64F

>>>> – `src.depth()` = CV_32F, ddepth = -1/CV_32F/CV_64F

>>>> – `src.depth()` = CV_64F, ddepth = -1/CV_64F

>>> when `ddepth=-1`, the destination image will have the same depth as the source.

>>> **kernel** – Convolution kernel (or rather a correlation kernel), a single-channel floating point matrix. If you want to apply different kernels to different channels, split the image into separate color planes using `split()` and process them individually.

>>> **anchor** – Anchor of the kernel that indicates the relative position of a filtered point within the kernel. The anchor should lie within the kernel. The special default value (-1,-1) means that the anchor is at the kernel center.

>>> **delta** – Optional value added to the filtered pixels before storing them in `dst` .

>>> **borderType** – Pixel extrapolation method. See `borderInterpolate()` for details.

The function applies an arbitrary linear filter to an image. In-place operation is supported. When the aperture is partially outside the image, the function interpolates outlier pixel values according to the specified border mode.

The function does actually compute correlation, not the convolution:

$$\mathtt{dst}(x,y) = \sum_{\substack{0 \le x' < \mathtt{kernel.cols}, \\ 0 \le y' < \mathtt{kernel.rows}}} \mathtt{kernel}(x',y') * \mathtt{src}(x + x' - \mathtt{anchor.x}, y + y' - \mathtt{anchor.y})$$

That is, the kernel is not mirrored around the anchor point. If you need a real convolution, flip the kernel using `flip()` and set the new anchor to (`kernel.cols - anchor.x - 1, kernel.rows - anchor.y - 1`) .

The function uses the DFT-based algorithm in case of sufficiently large kernels (~``11 x 11`` or larger) and the direct algorithm (that uses the engine retrieved by `createLinearFilter()` ) for small kernels.

**See Also:**

`sepFilter2D()`, `createLinearFilter()`, `dft()`, `matchTemplate()`

## GaussianBlur

Smoothes an image using a Gaussian filter.

**C++:** void **GaussianBlur**(InputArray **src**, OutputArray **dst**, Size **ksize**, double **sigmaX**, double **sigmaY**=0, int **borderType**=BORDER_DEFAULT )

**Python:** cv2.**GaussianBlur**(src, ksize, sigma1[, dst[, sigma2[, borderType ]]]) → dst

**Parameters**

**src** – Source image. The image can have any number of channels, which are processed independently. The depth should be `CV_8U`, `CV_16U`, `CV_16S`, `CV_32F` or `CV_64F`.

**dst** – Destination image of the same size and type as `src` .

**ksize** – Gaussian kernel size. `ksize.width` and `ksize.height` can differ but they both must be positive and odd. Or, they can be zero's and then they are computed from `sigma*` .

**sigmaX** – Gaussian kernel standard deviation in X direction.

**sigmaY** – Gaussian kernel standard deviation in Y direction. If `sigmaY` is zero, it is set to be equal to `sigmaX` . If both sigmas are zeros, they are computed from `ksize.width` and `ksize.height` , respectively. See `getGaussianKernel()` for details. To fully control the result regardless of possible future modifications of all this semantics, it is recommended to specify all of `ksize` , `sigmaX` , and `sigmaY` .

**borderType** – Pixel extrapolation method. See `borderInterpolate()` for details.

The function convolves the source image with the specified Gaussian kernel. In-place filtering is supported.

**See Also:**

`sepFilter2D()`, `filter2D()`, `blur()`, `boxFilter()`, `bilateralFilter()`, `medianBlur()`

## getDerivKernels

Returns filter coefficients for computing spatial image derivatives.

**C++:** void **getDerivKernels**(OutputArray **kx**, OutputArray **ky**, int **dx**, int **dy**, int **ksize**, bool **normalize**=false, int **ktype**=CV_32F )

**Python:** cv2.**getDerivKernels**(dx, dy, ksize[, kx[, ky[, normalize[, ktype ] ] ] ]) → kx, ky

**Parameters**

**kx** – Output matrix of row filter coefficients. It has the type `ktype` .

**ky** – Output matrix of column filter coefficients. It has the type `ktype` .

**dx** – Derivative order in respect of x.

**dy** – Derivative order in respect of y.

**ksize** – Aperture size. It can be `CV_SCHARR` , 1, 3, 5, or 7.

**normalize** – Flag indicating whether to normalize (scale down) the filter coefficients or not. Theoretically, the coefficients should have the denominator $= 2^{ksize*2-dx-dy-2}$ . If you are going to filter floating-point images, you are likely to use the normalized kernels. But if you compute derivatives of an 8-bit image, store the results in a 16-bit image, and wish to preserve all the fractional bits, you may want to set `normalize=false` .

**ktype** – Type of filter coefficients. It can be `CV_32f` or `CV_64F` .

The function computes and returns the filter coefficients for spatial image derivatives. When `ksize=CV_SCHARR` , the Scharr $3 \times 3$ kernels are generated (see `Scharr()` ). Otherwise, Sobel kernels are generated (see `Sobel()` ). The filters are normally passed to `sepFilter2D()` or to `createSeparableLinearFilter()` .

## getGaussianKernel

Returns Gaussian filter coefficients.

**C++:** Mat **getGaussianKernel**(int **ksize**, double **sigma**, int **ktype**=CV_64F )

**Python:** cv2.**getGaussianKernel**(ksize, sigma[, ktype]) → retval

> **Parameters**
>
>> **ksize** – Aperture size. It should be odd ( $\texttt{ksize} \mod 2 = 1$ ) and positive.
>>
>> **sigma** – Gaussian standard deviation. If it is non-positive, it is computed from `ksize` as `sigma = 0.3*((ksize-1)*0.5 - 1) + 0.8`.
>>
>> **ktype** – Type of filter coefficients. It can be `CV_32f` or `CV_64F` .

The function computes and returns the $\texttt{ksize} \times 1$ matrix of Gaussian filter coefficients:

$$G_i = \alpha * e^{-(i-(\texttt{ksize}-1)/2)^2/(2*\texttt{sigma})^2},$$

where $i = 0..\texttt{ksize} - 1$ and $\alpha$ is the scale factor chosen so that $\sum_i G_i = 1$.

Two of such generated kernels can be passed to `sepFilter2D()` or to `createSeparableLinearFilter()`. Those functions automatically recognize smoothing kernels (a symmetrical kernel with sum of weights equal to 1) and handle them accordingly. You may also use the higher-level `GaussianBlur()`.

**See Also:**

`sepFilter2D()`, `createSeparableLinearFilter()`, `getDerivKernels()`, `getStructuringElement()`, `GaussianBlur()`

## getKernelType

Returns the kernel type.

**C++:** int **getKernelType**(InputArray **kernel**, Point **anchor**)

> **Parameters**
>
>> **kernel** – 1D array of the kernel coefficients to analyze.
>>
>> **anchor** – Anchor position within the kernel.

The function analyzes the kernel coefficients and returns the corresponding kernel type:

- **KERNEL_GENERAL** The kernel is generic. It is used when there is no any type of symmetry or other properties.

- **KERNEL_SYMMETRICAL** The kernel is symmetrical: $\texttt{kernel}_i == \texttt{kernel}_{\texttt{ksize}-i-1}$ , and the anchor is at the center.

- **KERNEL_ASYMMETRICAL** The kernel is asymmetrical: $\texttt{kernel}_i == -\texttt{kernel}_{\texttt{ksize}-i-1}$ , and the anchor is at the center.

- **KERNEL_SMOOTH** All the kernel elements are non-negative and summed to 1. For example, the Gaussian kernel is both smooth kernel and symmetrical, so the function returns `KERNEL_SMOOTH | KERNEL_SYMMETRICAL` .

- **KERNEL_INTEGER** All the kernel coefficients are integer numbers. This flag can be combined with `KERNEL_SYMMETRICAL` or `KERNEL_ASYMMETRICAL` .

## getStructuringElement

Returns a structuring element of the specified size and shape for morphological operations.

**C++:** Mat **getStructuringElement**(int **shape**, Size **ksize**, Point **anchor**=Point(-1,-1))

**Python:** cv2.**getStructuringElement**(shape, ksize[, anchor]) → retval

**C:** IplConvKernel* **cvCreateStructuringElementEx**(int **cols**, int **rows**, int **anchorX**, int **anchorY**, int **shape**, int* **values**=NULL )

**Python:** cv.**CreateStructuringElementEx**(cols, rows, anchorX, anchorY, shape, values=None) → kernel

> **Parameters**
>
>> **shape** – Element shape that could be one of the following:
>>
>> – **MORPH_RECT** - a rectangular structuring element:
>>
>> $$E_{ij} = 1$$
>>
>> – **MORPH_ELLIPSE** - an elliptic structuring element, that is, a filled ellipse inscribed into the rectangle Rect(0, 0, esize.width, 0.esize.height)
>>
>> – **MORPH_CROSS** - a cross-shaped structuring element:
>>
>> $$E_{ij} = \begin{cases} 1 & \text{if i=anchor.y or j=anchor.x} \\ 0 & \text{otherwise} \end{cases}$$
>>
>> – **CV_SHAPE_CUSTOM** - custom structuring element (OpenCV 1.x API)
>>
>> **ksize** – Size of the structuring element.
>>
>> **cols** – Width of the structuring element
>>
>> **rows** – Height of the structuring element
>>
>> **anchor** – Anchor position within the element. The default value $(-1, -1)$ means that the anchor is at the center. Note that only the shape of a cross-shaped element depends on the anchor position. In other cases the anchor just regulates how much the result of the morphological operation is shifted.
>>
>> **anchorX** – x-coordinate of the anchor
>>
>> **anchorY** – y-coordinate of the anchor
>>
>> **values** – integer array of cols``*``rows elements that specifies the custom shape of the structuring element, when shape=CV_SHAPE_CUSTOM.

The function constructs and returns the structuring element that can be further passed to createMorphologyFilter(), erode(), dilate() or morphologyEx() . But you can also construct an arbitrary binary mask yourself and use it as the structuring element.

---

**Note:** When using OpenCV 1.x C API, the created structuring element IplConvKernel* element must be released in the end using cvReleaseStructuringElement(&element).

---

## medianBlur

Smoothes an image using the median filter.

**C++:** void **medianBlur**(InputArray **src**, OutputArray **dst**, int **ksize**)

---

**Python:** cv2.**medianBlur**(src, ksize[, dst ]) → dst

> **Parameters**
>
> > **src** – Source 1-, 3-, or 4-channel image. When `ksize` is 3 or 5, the image depth should be `CV_8U` , `CV_16U` , or `CV_32F` . For larger aperture sizes, it can only be `CV_8U` .
> >
> > **dst** – Destination array of the same size and type as `src` .
> >
> > **ksize** – Aperture linear size. It must be odd and greater than 1, for example: 3, 5, 7 ...

The function smoothes an image using the median filter with the `ksize` × `ksize` aperture. Each channel of a multi-channel image is processed independently. In-place operation is supported.

**See Also:**

bilateralFilter(), blur(), boxFilter(), GaussianBlur()

## morphologyEx

Performs advanced morphological transformations.

**C++:** void **morphologyEx**(InputArray **src**, OutputArray **dst**, int **op**, InputArray **element**, Point **anchor**=Point(-1,-1), int **iterations**=1, int **borderType**=BORDER_CONSTANT, const Scalar& **borderValue**=morphologyDefaultBorderValue() )

**Python:** cv2.**morphologyEx**(src, op, kernel[, dst[, anchor[, iterations[, borderType[, borderValue ] ] ] ] ]) → dst

**C:** void **cvMorphologyEx**(const CvArr* **src**, CvArr* **dst**, CvArr* **temp**, IplConvKernel* **element**, int **operation**, int **iterations**=1 )

**Python:** cv.**MorphologyEx**(src, dst, temp, element, operation, iterations=1) → None

> **Parameters**
>
> > **src** – Source image. The number of channels can be arbitrary. The depth should be one of `CV_8U, CV_16U, CV_16S, CV_32F'` or `''CV_64F`.
> >
> > **dst** – Destination image of the same size and type as `src` .
> >
> > **element** – Structuring element.
> >
> > **op** – Type of a morphological operation that can be one of the following:
> >
> > – **MORPH_OPEN** - an opening operation
> >
> > – **MORPH_CLOSE** - a closing operation
> >
> > – **MORPH_GRADIENT** - a morphological gradient
> >
> > – **MORPH_TOPHAT** - "top hat"
> >
> > – **MORPH_BLACKHAT** - "black hat"
> >
> > **iterations** – Number of times erosion and dilation are applied.
> >
> > **borderType** – Pixel extrapolation method. See borderInterpolate() for details.
> >
> > **borderValue** – Border value in case of a constant border. The default value has a special meaning. See createMorphologyFilter() for details.

The function can perform advanced morphological transformations using an erosion and dilation as basic operations.

Opening operation:

$$dst = \text{open}(src, \text{element}) = \text{dilate}(\text{erode}(src, \text{element}))$$

Closing operation:

$$dst = close(\mathsf{src}, \mathsf{element}) = erode(dilate(\mathsf{src}, \mathsf{element}))$$

Morphological gradient:

$$dst = morph\_grad(\mathsf{src}, \mathsf{element}) = dilate(\mathsf{src}, \mathsf{element}) - erode(\mathsf{src}, \mathsf{element})$$

"Top hat":

$$dst = tophat(\mathsf{src}, \mathsf{element}) = \mathsf{src} - open(\mathsf{src}, \mathsf{element})$$

"Black hat":

$$dst = blackhat(\mathsf{src}, \mathsf{element}) = close(\mathsf{src}, \mathsf{element}) - \mathsf{src}$$

Any of the operations can be done in-place. In case of multi-channel images, each channel is processed independently.

**See Also:**

dilate(), erode(), createMorphologyFilter()

## Laplacian

Calculates the Laplacian of an image.

**C++:** void **Laplacian**(InputArray **src**, OutputArray **dst**, int **ddepth**, int **ksize**=1, double **scale**=1, double **delta**=0, int **borderType**=BORDER_DEFAULT )

**Python:** cv2.**Laplacian**(src, ddepth[, dst[, ksize[, scale[, delta[, borderType ]]]]]) $\rightarrow$ dst

**C:** void **cvLaplace**(const CvArr* **src**, CvArr* **dst**, int **ksize**=3)

**Python:** cv.**Laplace**(src, dst, ksize=3) $\rightarrow$ None

>       **Parameters**

>> **src** – Source image.

>> **dst** – Destination image of the same size and the same number of channels as src .

>> **ddepth** – Desired depth of the destination image.

>> **ksize** – Aperture size used to compute the second-derivative filters. See getDerivKernels() for details. The size must be positive and odd.

>> **scale** – Optional scale factor for the computed Laplacian values. By default, no scaling is applied. See getDerivKernels() for details.

>> **delta** – Optional delta value that is added to the results prior to storing them in dst .

>> **borderType** – Pixel extrapolation method. See borderInterpolate() for details.

The function calculates the Laplacian of the source image by adding up the second x and y derivatives calculated using the Sobel operator:

$$dst = \Delta \mathsf{src} = \frac{\partial^2 \mathsf{src}}{\partial x^2} + \frac{\partial^2 \mathsf{src}}{\partial y^2}$$

This is done when ksize > 1 . When ksize == 1 , the Laplacian is computed by filtering the image with the following $3 \times 3$ aperture:

$$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$$

**See Also:**

Sobel(), Scharr()

## pyrDown

Smoothes an image and downsamples it.

**C++:** void **pyrDown**(InputArray **src**, OutputArray **dst**, const Size& **dstsize**=Size())

**Python:** cv2.**pyrDown**(src[, dst[, dstsize ]]) → dst

**C:** void **cvPyrDown**(const CvArr* **src**, CvArr* **dst**, int **filter**=CV_GAUSSIAN_5x5 )

**Python:** cv.**PyrDown**(src, dst, filter=CV_GAUSSIAN_5X5) → None

> **Parameters**

>> **src** – Source image.

>> **dst** – Destination image. It has the specified size and the same type as src .

>> **dstsize** – Size of the destination image. By default, it is computed as Size((src.cols+1)/2, (src.rows+1)/2) . But in any case, the following conditions should be satisfied:

$$|\texttt{dstsize.width} * 2 - \texttt{src.cols}| \leq 2$$
$$|\texttt{dstsize.height} * 2 - \texttt{src.rows}| \leq 2$$

The function performs the downsampling step of the Gaussian pyramid construction. First, it convolves the source image with the kernel:

$$\frac{1}{256}\begin{bmatrix} 1 & 4 & 6 & 4 & 1 \\ 4 & 16 & 24 & 16 & 4 \\ 6 & 24 & 36 & 24 & 6 \\ 4 & 16 & 24 & 16 & 4 \\ 1 & 4 & 6 & 4 & 1 \end{bmatrix}$$

Then, it downsamples the image by rejecting even rows and columns.

## pyrUp

Upsamples an image and then smoothes it.

**C++:** void **pyrUp**(InputArray **src**, OutputArray **dst**, const Size& **dstsize**=Size())

**Python:** cv2.**pyrUp**(src[, dst[, dstsize ]]) → dst

**C:** **cvPyrUp**(const CvArr* **src**, CvArr* **dst**, int **filter**=CV_GAUSSIAN_5x5 )

**Python:** cv.**PyrUp**(src, dst, filter=CV_GAUSSIAN_5X5) → None

> **Parameters**

>> **src** – Source image.

>> **dst** – Destination image. It has the specified size and the same type as src .

>> **dstsize** – Size of the destination image. By default, it is computed as Size(src.cols*2, (src.rows*2) . But in any case, the following conditions should be satisfied:

$$|\texttt{dstsize.width} - \texttt{src.cols} * 2| \leq (\texttt{dstsize.width} \mod 2)$$
$$|\texttt{dstsize.height} - \texttt{src.rows} * 2| \leq (\texttt{dstsize.height} \mod 2)$$

The function performs the upsampling step of the Gaussian pyramid construction though it can actually be used to construct the Laplacian pyramid. First, it upsamples the source image by injecting even zero rows and columns and then convolves the result with the same kernel as in `pyrDown()` multiplied by 4.

## pyrMeanShiftFiltering

Performs initial step of meanshift segmentation of an image.

**C++:** void **pyrMeanShiftFiltering**(InputArray **src**, OutputArray **dst**, double **sp**, double **sr**, int **maxLevel**=1, TermCriteria **termcrit**=TermCriteria(TermCriteria::MAX_ITER+TermCriteria::EPS,5,1) )

**Python:** cv2.**pyrMeanShiftFiltering**(src, sp, sr[, dst[, maxLevel[, termcrit]]]) → dst

**C:** void **cvPyrMeanShiftFiltering**(const CvArr* **src**, CvArr* **dst**, double **sp**, double **sr**, int **max_level**=1, CvTermCriteria **termcrit**=cvTermCriteria(CV_TERMCRIT_ITER+CV_TERMCRIT_EPS,5,1))

**Python:** cv.**PyrMeanShiftFiltering**(src, dst, sp, sr, maxLevel=1, termcrit=(CV_TERMCRIT_ITER+CV_TERMCRIT_EPS, 5, 1)) → None

> **Parameters**
>
> > **src** – The source 8-bit, 3-channel image.
> >
> > **dst** – The destination image of the same format and the same size as the source.
> >
> > **sp** – The spatial window radius.
> >
> > **sr** – The color window radius.
> >
> > **maxLevel** – Maximum level of the pyramid for the segmentation.
> >
> > **termcrit** – Termination criteria: when to stop meanshift iterations.

The function implements the filtering stage of meanshift segmentation, that is, the output of the function is the filtered "posterized" image with color gradients and fine-grain texture flattened. At every pixel (`X,Y`) of the input image (or down-sized input image, see below) the function executes meanshift iterations, that is, the pixel (`X,Y`) neighborhood in the joint space-color hyperspace is considered:

$$(x, y) : X - \mathsf{sp} \le x \le X + \mathsf{sp}, Y - \mathsf{sp} \le y \le Y + \mathsf{sp}, \|(R, G, B) - (r, g, b)\| \le \mathsf{sr}$$

where (`R,G,B`) and (`r,g,b`) are the vectors of color components at (`X,Y`) and (`x,y`), respectively (though, the algorithm does not depend on the color space used, so any 3-component color space can be used instead). Over the neighborhood the average spatial value (`X',Y'`) and average color vector (`R',G',B'`) are found and they act as the neighborhood center on the next iteration:

$$(X, Y) \ (X', Y'), (R, G, B) \ (R', G', B').$$

After the iterations over, the color components of the initial pixel (that is, the pixel from where the iterations started) are set to the final value (average color at the last iteration):

$$I(X, Y) < -(R*, G*, B*)$$

When `maxLevel > 0`, the gaussian pyramid of `maxLevel+1` levels is built, and the above procedure is run on the smallest layer first. After that, the results are propagated to the larger layer and the iterations are run again only on those pixels where the layer colors differ by more than `sr` from the lower-resolution layer of the pyramid. That makes boundaries of color regions sharper. Note that the results will be actually different from the ones obtained by running the meanshift procedure on the whole original image (i.e. when `maxLevel==0`).

## sepFilter2D

Applies a separable linear filter to an image.

**C++:** void **sepFilter2D**(InputArray **src**, OutputArray **dst**, int **ddepth**, InputArray **rowKernel**, InputArray **columnKernel**, Point **anchor**=Point(-1,-1), double **delta**=0, int **borderType**=BORDER_DEFAULT )

**Python:** cv2.**sepFilter2D**(src, ddepth, kernelX, kernelY[, dst[, anchor[, delta[, borderType ]]]]) → dst

> **Parameters**
>
> > **src** – Source image.
> >
> > **dst** – Destination image of the same size and the same number of channels as `src` .
> >
> > **ddepth** –
> >
> > **Destination image depth. The following combination of `src.depth()` and `ddepth` are supported:**
> >
> > > – `src.depth()` = CV_8U, ddepth = -1/CV_16S/CV_32F/CV_64F
> > >
> > > – `src.depth()` = CV_16U/CV_16S, ddepth = -1/CV_32F/CV_64F
> > >
> > > – `src.depth()` = CV_32F, ddepth = -1/CV_32F/CV_64F
> > >
> > > – `src.depth()` = CV_64F, ddepth = -1/CV_64F
> >
> > when `ddepth=-1`, the destination image will have the same depth as the source.
> >
> > **rowKernel** – Coefficients for filtering each row.
> >
> > **columnKernel** – Coefficients for filtering each column.
> >
> > **anchor** – Anchor position within the kernel. The default value $(-1, 1)$ means that the anchor is at the kernel center.
> >
> > **delta** – Value added to the filtered results before storing them.
> >
> > **borderType** – Pixel extrapolation method. See `borderInterpolate()` for details.

The function applies a separable linear filter to the image. That is, first, every row of `src` is filtered with the 1D kernel `rowKernel` . Then, every column of the result is filtered with the 1D kernel `columnKernel` . The final result shifted by `delta` is stored in `dst` .

**See Also:**

`createSeparableLinearFilter()`, `filter2D()`, `Sobel()`, `GaussianBlur()`, `boxFilter()`, `blur()`

## Smooth

Smooths the image in one of several ways.

**C:** void **cvSmooth**(const CvArr* **src**, CvArr* **dst**, int **smoothtype**=CV_GAUSSIAN, int **param1**=3, int **param2**=0, double **param3**=0, double **param4**=0)

**Python:** cv.**Smooth**(src, dst, smoothtype=CV_GAUSSIAN, param1=3, param2=0, param3=0, param4=0) → None

> **Parameters**
>
> > **src** – The source image
> >
> > **dst** – The destination image
> >
> > **smoothtype** – Type of the smoothing:
> >
> > – **CV_BLUR_NO_SCALE** linear convolution with `param1` × `param2` box kernel (all 1's). If you want to smooth different pixels with different-size box kernels, you can use the integral image that is computed using `integral()`
> >
> > – **CV_BLUR** linear convolution with `param1`×`param2` box kernel (all 1's) with subsequent scaling by $1/(\texttt{param1} \cdot \texttt{param2})$
> >
> > – **CV_GAUSSIAN** linear convolution with a `param1` × `param2` Gaussian kernel
> >
> > – **CV_MEDIAN** median filter with a `param1` × `param1` square aperture
> >
> > – **CV_BILATERAL** bilateral filter with a `param1` × `param1` square aperture, color sigma= `param3` and spatial sigma= `param4`. If `param1=0`, the aperture square side is set to `cvRound(param4*1.5)*2+1`. Information about bilateral filtering can be found at http://www.dai.ed.ac.uk/CVonline/LOCAL_COPIES/MANDUCHI1/Bilateral_Filtering.html
> >
> > **param1** – The first parameter of the smoothing operation, the aperture width. Must be a positive odd number (1, 3, 5, ...)
> >
> > **param2** – The second parameter of the smoothing operation, the aperture height. Ignored by `CV_MEDIAN` and `CV_BILATERAL` methods. In the case of simple scaled/non-scaled and Gaussian blur if `param2` is zero, it is set to `param1`. Otherwise it must be a positive odd number.
> >
> > **param3** – In the case of a Gaussian parameter this parameter may specify Gaussian σ (standard deviation). If it is zero, it is calculated from the kernel size:
> >
> > $$\sigma = 0.3(n/2 - 1) + 0.8 \quad \text{where} \quad n = \begin{array}{l} \texttt{param1 for horizontal kernel} \\ \texttt{param2 for vertical kernel} \end{array}$$
> >
> > Using standard sigma for small kernels ( $3 \times 3$ to $7 \times 7$ ) gives better speed. If `param3` is not zero, while `param1` and `param2` are zeros, the kernel size is calculated from the sigma (to provide accurate enough operation).

The function smooths an image using one of several methods. Every of the methods has some features and restrictions listed below:

- Blur with no scaling works with single-channel images only and supports accumulation of 8-bit to 16-bit format (similar to `Sobel()` and `Laplacian()`) and 32-bit floating point to 32-bit floating-point format.

- Simple blur and Gaussian blur support 1- or 3-channel, 8-bit and 32-bit floating point images. These two methods can process images in-place.

- Median and bilateral filters work with 1- or 3-channel 8-bit images and can not process images in-place.

**Note:** The function is now obsolete. Use `GaussianBlur()`, `blur()`, `medianBlur()` or `bilateralFilter()`.

## Sobel

Calculates the first, second, third, or mixed image derivatives using an extended Sobel operator.

**C++:** void **Sobel**(InputArray **src**, OutputArray **dst**, int **ddepth**, int **xorder**, int **yorder**, int **ksize**=3, double **scale**=1, double **delta**=0, int **borderType**=BORDER_DEFAULT )

**Python:** cv2.**Sobel**(src, ddepth, dx, dy[, dst[, ksize[, scale[, delta[, borderType ]]]]]) → dst

**C:** void **cvSobel**(const CvArr* **src**, CvArr* **dst**, int **xorder**, int **yorder**, int **apertureSize**=3 )

**Python:** cv.**Sobel**(src, dst, xorder, yorder, apertureSize=3) → None

> **Parameters**

>> **src** – Source image.

>> **dst** – Destination image of the same size and the same number of channels as `src` .

>> **ddepth** –

>> **Destination image depth. The following combination of `src.depth()` and `ddepth` are supported:**

>>> – `src.depth()` = CV_8U, ddepth = -1/CV_16S/CV_32F/CV_64F

>>> – `src.depth()` = CV_16U/CV_16S, ddepth = -1/CV_32F/CV_64F

>>> – `src.depth()` = CV_32F, ddepth = -1/CV_32F/CV_64F

>>> – `src.depth()` = CV_64F, ddepth = -1/CV_64F

>> when `ddepth=-1`, the destination image will have the same depth as the source. In the case of 8-bit input images it will result in truncated derivatives.

>> **xorder** – Order of the derivative x.

>> **yorder** – Order of the derivative y.

>> **ksize** – Size of the extended Sobel kernel. It must be 1, 3, 5, or 7.

>> **scale** – Optional scale factor for the computed derivative values. By default, no scaling is applied. See `getDerivKernels()` for details.

>> **delta** – Optional delta value that is added to the results prior to storing them in `dst` .

>> **borderType** – Pixel extrapolation method. See `borderInterpolate()` for details.

In all cases except one, the $\mathtt{ksize} \times \mathtt{ksize}$ separable kernel is used to calculate the derivative. When $\mathtt{ksize} = 1$, the $3 \times 1$ or $1 \times 3$ kernel is used (that is, no Gaussian smoothing is done). $\mathtt{ksize} = 1$ can only be used for the first or the second x- or y- derivatives.

There is also the special value $\mathtt{ksize} = \mathtt{CV\_SCHARR}$ (-1) that corresponds to the $3 \times 3$ Scharr filter that may give more accurate results than the $3 \times 3$ Sobel. The Scharr aperture is

$$\begin{bmatrix} -3 & 0 & 3 \\ -10 & 0 & 10 \\ -3 & 0 & 3 \end{bmatrix}$$

for the x-derivative, or transposed for the y-derivative.

The function calculates an image derivative by convolving the image with the appropriate kernel:

$$\mathtt{dst} = \frac{\partial^{\mathtt{xorder+yorder}}\mathtt{src}}{\partial x^{\mathtt{xorder}}\partial y^{\mathtt{yorder}}}$$

The Sobel operators combine Gaussian smoothing and differentiation, so the result is more or less resistant to the noise. Most often, the function is called with ( $\mathtt{xorder} = 1$, $\mathtt{yorder} = 0$, $\mathtt{ksize} = 3$) or ( $\mathtt{xorder} = 0$, $\mathtt{yorder} = 1$,

`ksize` = 3) to calculate the first x- or y- image derivative. The first case corresponds to a kernel of:

$$\begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

The second case corresponds to a kernel of:

$$\begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$

**See Also:**

Scharr(), Laplacian(), sepFilter2D(), filter2D(), GaussianBlur(), cartToPolar()

## Scharr

Calculates the first x- or y- image derivative using Scharr operator.

**C++:** void **Scharr**(InputArray **src**, OutputArray **dst**, int **ddepth**, int **xorder**, int **yorder**, double **scale**=1, double **delta**=0, int **borderType**=BORDER_DEFAULT )

**Python:** cv2.**Scharr**(src, ddepth, dx, dy[, dst[, scale[, delta[, borderType ]]]]) → dst

>       **Parameters**
>
>>           **src** – Source image.
>>
>>           **dst** – Destination image of the same size and the same number of channels as `src`.
>>
>>           **ddepth** – Destination image depth. See Sobel() for the list of supported combination of `src.depth()` and `ddepth`.
>>
>>           **xorder** – Order of the derivative x.
>>
>>           **yorder** – Order of the derivative y.
>>
>>           **scale** – Optional scale factor for the computed derivative values. By default, no scaling is applied. See getDerivKernels() for details.
>>
>>           **delta** – Optional delta value that is added to the results prior to storing them in `dst`.
>>
>>           **borderType** – Pixel extrapolation method. See borderInterpolate() for details.

The function computes the first x- or y- spatial image derivative using the Scharr operator. The call

>       Scharr(src, dst, ddepth, xorder, yorder, scale, delta, borderType)

is equivalent to

>       Sobel(src, dst, ddepth, xorder, yorder, CV_SCHARR, scale, delta, borderType).

**See Also:**

cartToPolar()